Haokun Zhu

+1 4129548608 haokunz@andrew.cmu.edu haokunzhu.com Haokun Zhu zwandering

Education

- **Carnegie Mellon University** — Pittsburgh, United States
  Ph.D. in Robotics; Advisor: Ji Zhang — *Aug. 2026 – Present*
- **Carnegie Mellon University** — Pittsburgh, United States
  Master of Science in Robotics; **GPA: 4.25/4.3** — *Aug. 2024 – Jun. 2026*
- **Shanghai Jiao Tong University** — Shanghai, China
  B.Eng in Computer Science and Technology; **GPA: 89.14/100; 3.81/4.3** — *Sep. 2020 – Jun. 2024*

Publications (* means equal contribution, ♯ means corresponding author)

- **SysNav: Multi-Level Systematic Cooperation Enables Real-World, Cross-Embodiment Object Navigation**
  Haokun Zhu, Zongtai Li, Zihan Liu, Kevin Guo, Zhengzhi Lin, Yuxin Cai, Guofei Chen, Chen Lv, Wenshan Wang, Jean Oh, Ji Zhang
  IEEE RA-L 2026. [**Website**] [**ArXiv:2603.06914**]

- **STRIVE: Structured Representation Integrating VLM Reasoning for Efficient Object Navigation**
  Haokun Zhu*, Zongtai Li*, Zhixuan Liu, Wenshan Wang, Ji Zhang, Jonathan Francis, Jean Oh
  ICRA 2026; RSS 2025 SemRob Workshop. [**Website**] [**ArXiv:2505.06729**]

- **MOSAIC: Generating Consistent, Privacy-Preserving Scenes from Multiple Depth Views in Multi-Room Environments**
  Zhixuan Liu, Haokun Zhu, Rui Chen, Jonathan Francis, Soonmin Hwang, Ji Zhang, Jean Oh
  ICCV 2025. [**Website**] [**ArXiv:2503.13816**]

- **AesStyler: Aesthetic Guided Universal Style Transfer**
  Ran Yi*,♯, Haokun Zhu*, Teng Hu, Yu-Kun Lai, Paul L. Rosin
  ACM MM 2024. [**PDF**]

- **SAMVG: A Multi-stage Image Vectorization Model with the Segment-Anything Model**
  Haokun Zhu*, Juang Ian Chong*, Teng Hu, Ran Yi♯, Yu-Kun Lai, Paul L. Rosin
  ICASSP 2024. [**ArXiv:2311.05276**]

- **M3DM-NR: RGB-3D Noisy-Resistant Industrial Anomaly Detection via Multimodal Denoising**
  Chengjie Wang, Haokun Zhu, Jinlong Peng, Yue Wang, Ran Yi, Yunsheng Wu, Lizhuang Ma, Jiangning Zhang
  TPAMI 2025. [**ArXiv:2406.02263**]

- **Stroke-based Neural Painting and Stylization with Dynamically Predicted Painting Region**
  Teng Hu, Ran Yi♯, Haokun Zhu, Liang Liu, Jinlong Peng, Yabiao Wang, Chengjie Wang, Lizhuang Ma
  ACM MM 2023. [**Code**] [**ArXiv:2309.03504**]

- **Phasic Content Fusing Diffusion Model with Directional Distribution Consistency for Few-Shot Model Adaption**
  Teng Hu, Jiangning Zhang, Liang Liu, Ran Yi♯, Siqi Kou, Haokun Zhu, Xu Chen, Yabiao Wang, Chengjie Wang Lizhuang Ma
  ICCV 2023. [**Code**] [**ArXiv:2309.03729**]

Experience & Internship

- **Zhang Lab in CMU** — Pittsburgh, United States
  Ph.D. Student advised by Dr. Ji Zhang — *Aug. 2026 – Present*
  - **JEPA World Model for Navigation**: exploring how an action-conditioned V-JEPA world model can serve navigation, especially tasks that require understanding other agents' intentions.
- **roBot Intelligence Group (BIG) in CMU** — Pittsburgh, United States
  Graduate Research Assistant co-advised by Prof. Jean Oh and Dr. Ji Zhang — *Sep. 2024 – Jun. 2026*
  - **VLM-Guided Object Navigation in Simulation**: designed structured multi-layer scene representation (viewpoint, object, and room nodes) and two-stage navigation policy combining VLM reasoning with classical exploration methods, achieving state-of-the-art performance across HM3D, RoboTHOR, and MP3D benchmarks (ICRA 2026).
  - **SysNav: Cross-Embodiment Navigation System in Real-World**: built a three-level system for building-scale object navigation, deployed on wheeled, quadruped, and humanoid robots across 190 real-world experiments with 4-5x higher efficiency than baselines (RA-L 2026).
- **Digital Media Computer Vision Laboratory(DMCV) in SJTU** — Shanghai, China
  Undergraduate Research Assistant advised by Prof Ran Yi — *Oct. 2022 – Jun. 2024*

- **Image Generation & Stylization**: few-shot diffusion model adaption (ICCV 2023), aesthetic guided style transfer (ACM MM 2024), and SAM-based image vectorization (ICASSP 2024).

- **Youtu Lab at Tencent Technology (Shanghai) Co.Ltd** Shanghai, China
  Research Intern advised by Jinlong Peng *Dec. 2023 – Jun. 2024*
  - **Multimodal Industrial Anomaly Detection**: addressed ineffective feature integration issues in 3D point cloud and RGB images by applying multimodality to enhance industrial anomaly detection.

Research Projects

- **JEPA World Model for Navigation (Robotics, World Models)** *Jul. 2026 - Present*
  Advisor: Ji Zhang
  We study how an action-conditioned JEPA world model can serve navigation, especially in dynamic scenes where the robot must anticipate the intentions of other agents (e.g., pedestrians). In this project,
  - **I built action-conditioned world models in the frozen V-JEPA 2.1 latent space that imagine how a scene, including other agents, evolves under candidate robot actions.**
  - **I compared regression and flow-matching predictors with camera-ray conditioning, trained on 710 hours of real-world robot data.**
  - **I am developing world-model-based planning that reasons about other agents' intentions.**

- **SysNav: Cross-Embodiment Object Navigation (Robotics)** *May. 2025 - Jun. 2026*
  Advisor: Ji Zhang, Jean Oh
  We propose SysNav, a three-level object navigation system that decouples semantic reasoning, navigation planning, and motion control, treating rooms as minimal decision-making units and supporting both standard and conditional object navigation in building-scale real-world environments. In this project,
  - **I integrated VLM reasoning into the ROS2-based TARE Planner and deployed the system on a wheeled robot, Unitree Go2 quadruped, and Unitree G1 humanoid, validating the embodiment-agnostic design.**
  - **I evaluated the system through 190 real-world experiments, achieving 4-5x higher navigation efficiency than existing baselines.**
  - **This work is accepted by IEEE Robotics and Automation Letters (RA-L) 2026.**

- **STRIVE: VLM-Guided Object Navigation in Simulation (Robotics)** *Sep. 2024 - May. 2025*
  Advisor: Jean Oh, Ji Zhang, Jonathan Francis
  We propose STRIVE, a framework that enables robots to incrementally build a multi-layer scene representation during exploration and leverages VLM reasoning for efficient object navigation. We designed a two-stage navigation policy combining high-level VLM-guided planning with low-level exploration methods. In this project,
  - **I led the design of the structured scene representation (viewpoint, object, and room nodes) and implemented the two-stage VLM-guided navigation policy.**
  - **I conducted extensive benchmark evaluations across HM3D v1/v2, RoboTHOR, and MP3D, achieving state-of-the-art performance in both success rate (↑7.1%) and efficiency (↑12.5%).**
  - **This work is accepted by ICRA 2026 and the RSS 2025 SemRob Workshop.**

- **AesStyler: Aesthetic Guided Universal Style Transfer (CV)** *Jun. 2023 - Dec. 2023*
  Advisor: Ran Yi, Yu-Kun Lai, Paul L. Rosin
  We propose AesStyler, a novel Aesthetic Guided Universal Style Transfer method, utilizing pre-trained aesthetic assessment model, a novel Universal Aesthetic Codebook and a novel Universal and Specific Aesthetic-Guided Attention module. Extensive experiments and user-studies have demonstrated that our approach generates aesthetically more harmonious and pleasing results than the state-of-the-art methods. In this project,
  - **I proposed and independently implemented three novel contributions, and drafted the initial paper including methodology, figures, and result analysis.**
  - **I single-handedly conducted all qualitative and quantitative experiments and 2 user-studies.**
  - **This work is accepted by ACM MM 2024.**

- **SAMVG: An Image Vectorization Model with SAM (CV, CG)** *Aug. 2023 - Dec. 2023*
  Advisor: Ran Yi, Yu-Kun Lai, Paul L. Rosin
  We propose SAMVG, a multi-stage model to vectorize raster images into Scalable Vector Graphics. Extensive experiments demonstrate that SAMVG can produce high quality SVGs in any domain with less computation time and complexity compared to previous SOTA methods. In this project,
  - **I co-proposed three innovative aspects with the co-author and finished the final draft of the paper, including the methodology and result analysis.**
  - **This work is accepted by ICASSP 2024.**

- **M3DM-NR: RGB-3D Noisy Industrial Anomaly Detection (CV)** *Dec. 2023 - Jun. 2024*
  Advisor: Jinlong Peng
  We propose M3DM-NR, a noise-resistant framework leveraging the multi-modal (image and point cloud) discriminative capabilities of CLIP for noisy 3D-RGB anomaly detection. In this project,
  - **I proposed the 3-stage noise-resistant framework, designed and conducted all qualitative and quantitative experiments, and drafted the initial paper.**
  - **This work is accepted by TPAMI.**

Technical Skills

- **Languages** Python, C/C++, Matlab, LaTex, Mandarin(native), English(fluent, TOEFL: 111)
- **Tools** ROS, PyTorch, TensorFlow, OpenCV, OpenGL, LaTex, Markdown, git